

CS1217 - Spring 2023 - Homework 2

Bhumika Mittal, Saptarishi Dhanuka

Question 1

proc is a pseudo-filesystem which provides an interface to kernel data structures. Most of its files are read-only as it is just a virtual filesystem. It contains information about the processes that are currently running.

```
cs1217@cs1217-devel:~$ cd /proc
cs1217@cs1217-devel:/proc$ ls
1      107  127  1516 16    1688 1780 1841 1899 198    2140 2240 264 294 33    614 676 791 87 99    devices      ioports      locks         schedstat     timer_list
10     1082 1280 1528 1822 169    1788 1842 19    1993    2148 2242 27    295 337 617 677 8    88    acpi          diskstats    lra           mdstat        scsi
100    109    13    1532 1634 1699 1792 1848 1900 1998    2150 23    28    296 37    615 682 88 89    dma           kallsyms     meninfo       self           tty
101    11    14    1537 1653 17    18    1849 1905 2    2172 2314 280 297 4    619 691 81 9    bootconfig   driver        kcore         misc           slabinfo      version
102    110  1434 1541 166    1701 1816 1854 1907 20    2173 2379 288 298 478 623 699 82 90    buddyinfo    dynamic_debug keys          modules        softirqs      version_signature
103    1110 15    1549 1669 1705 1818 1855 1918 2020    2182 238 289 3    481 631 7    83 92    bus          execdomains  key-users     mounts         stat           vmallocInfo
104    118    1500 1553 187    1737 1826 1856 1919 2046    2184 24    29 30 5    635 719 84 93    cgroups      fb           knsg          ntrr           swaps          vmsstat
105    12    1507 1560 1678 175    1829 1861 1924 2077    2186 2437 280 300 570 637 722 847 95    cmdline      filesystems  kpagecgroup   nrt            sys            zoneinfo
106    121  1508 1561 168    1759 1833 1863 1962 2081    2187 25    291 31 6    644 744 849 96    consoles     fs           kpagecount    pagetypeinfo  sysrq-trigger
1062   122  1514 1565 1683 1765 1834 1867 197 21 22 26 292 32 612 663 747 85 97    cpuinfo      interrupts  kpageflags    partitions     sysvipc
1065   1244 1515 1564 1688    1773 1836 1880 1971 2108    2212 263 293 329 613 672 749 86 98    crypto       iomem       loadavg       pressure       thread-self
```

Question 2

The following command-line arguments can be provided to get only the processes that belong to a specified user, say cs304:

(a) **top -U cs304**

top provides the dynamic real-time view of the processes that are *running currently*. The -U is used to sort the processes that belong to the user cs304.

(b) **ps -u cs304**

ps stands for process status. Like top, it provides the view of all the processes. -u is used for displaying the processes for the given effective user ID (which is like a virtual ID on the basis of which OS makes decisions).

(c) **ps -U cs304**

-U is used for displaying the processes for the given real user ID (real UID is the ID of the user which owns the process).

Question 3

(a) The core requirements of any shell are:

- User interface which allows the user to interact with the operating system either through textual commands in a Command-Line Interpreter or through direct manipulation of graphical elements in a GUI. In a CLI, it needs to support commands which could have many arguments
- Have functionality for Input/Output redirection from and to files and commands
- Have functionality to implement piping of output of a command as input of another command
- Have builtin commands which don't have to be exec'ed by the shell separately such as `cd` and `echo` in `bash`.
- Storing environment variables like `$PATH` and `$HOME`
- Support for traversing the file system by changing directories

(b) The OS mainly needs to provide support to the shell by having a system call interface which actually implements functionality related to `fork()`, `exec()`, opening and closing files, `pipe()`, managing processes, doing I/O and anything else that can only be done in the privileged kernel mode. It also needs to have the support for environment variables such as `$PATH`

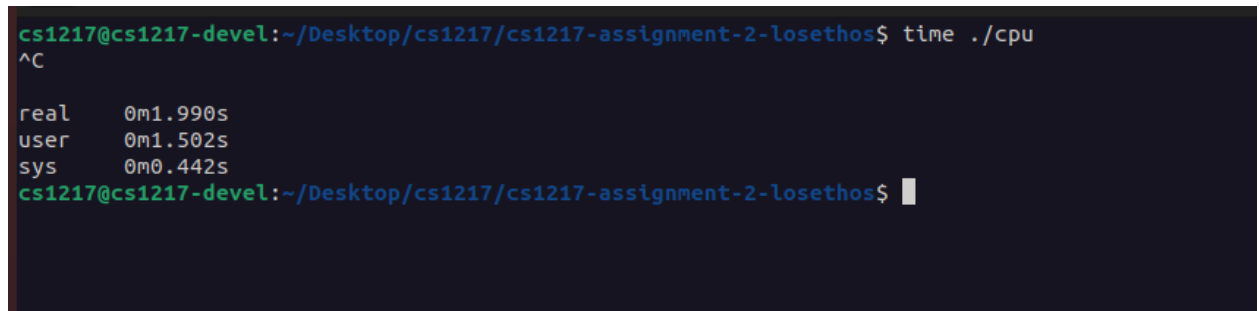
(c) The following steps will need to be implemented to have a basic, working shell program:

- Check if the 3 main file descriptors associated with `stdin`, `stdout` and `stderr` are open
- Have a main loop which reads a line of input from the user
- Parse the user input into the command and its arguments
- Start a process by calling the `fork()` syscall which duplicates the process into a child process.
- If the return value of `fork()` is 0 then we are in the child process and we call the `exec()` syscall with an argument which is the path to an executable binary file. This transforms the child process into a process specified by the binary like `ls`, `ps`, etc.
- If the return value of `fork()` is nonzero then we wait for the child process to finish

Question 4

(a) **cpu:**

Using the time command in linux, we can observe that this program spends most of its time (around 80%) in the user mode whereas it spends the rest 20% time in the kernel mode. In the following screenshot, the cpu program was run for 1.9 seconds before it was terminated using SIGKILL.



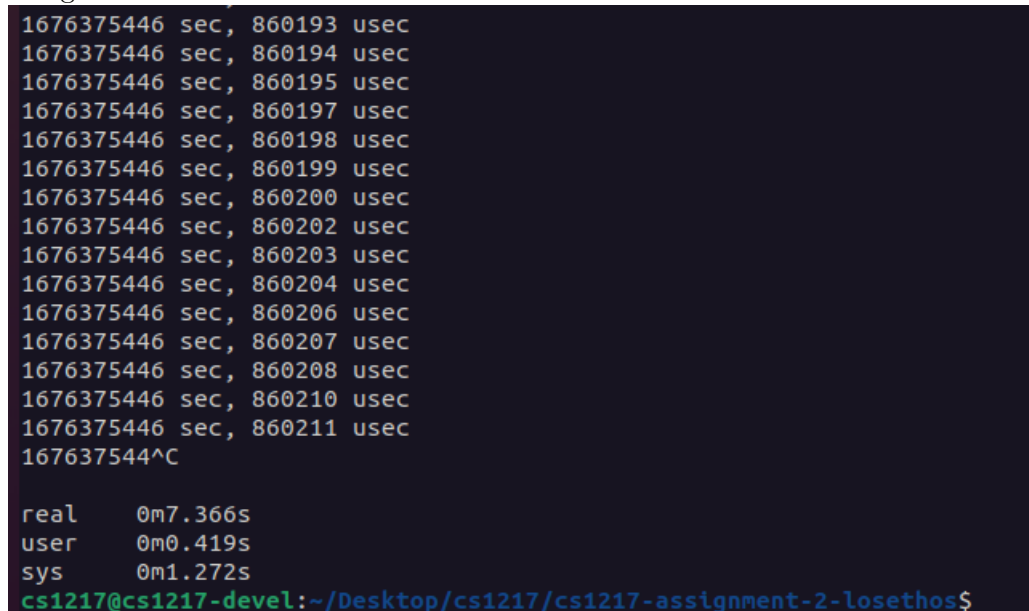
```
cs1217@cs1217-devel:~/Desktop/cs1217/cs1217-assignment-2-losethos$ time ./cpu
^C

real    0m1.990s
user    0m1.502s
sys     0m0.442s
cs1217@cs1217-devel:~/Desktop/cs1217/cs1217-assignment-2-losethos$
```

From the code, we can observe that the program spends most of its time doing calculations using CPU, hence it stays in the user mode. It goes to kernel mode initially while importing libraries, assigning values to the variables and over-writing the computed values.

(b) **cpu-print:**

Using the time command in linux, we can observe that this program spends most of its time (around 90%) in the kernel mode whereas it spends the rest 10% time in the user mode. In the following screenshot, the cpu-print program was run for 7.36 seconds before it was terminated using SIGKILL.



```
1676375446 sec, 860193 usec
1676375446 sec, 860194 usec
1676375446 sec, 860195 usec
1676375446 sec, 860197 usec
1676375446 sec, 860198 usec
1676375446 sec, 860199 usec
1676375446 sec, 860200 usec
1676375446 sec, 860202 usec
1676375446 sec, 860203 usec
1676375446 sec, 860204 usec
1676375446 sec, 860206 usec
1676375446 sec, 860207 usec
1676375446 sec, 860208 usec
1676375446 sec, 860210 usec
1676375446 sec, 860211 usec
1676375446^C

real    0m7.366s
user    0m0.419s
sys     0m1.272s
cs1217@cs1217-devel:~/Desktop/cs1217/cs1217-assignment-2-losethos$
```

From the code, we can observe that the program spends most of its time doing two functions (gettimeofday() and printf()) which in turn call syscalls which are used to get the time of the day and then print respectively. The computation for incrementing the count value, etc is done in the user mode.

Question 5

- (a) The shell program is **bash**.
- (b) It is maintained the variable **SHELL**.
- (c) To print the same, we can use **echo \$SHELL**.
- (d) The pid of the shell could be printed using **echo \$\$**, which shows the pid of the current shell as 2314.
- (e) The process tree is as follows:
systemd(1)---systemd(1507)---gnome-terminal-(2240)---bash(2314)---pstree(2699)

Here, systemd(1) is init, since it's pid is 1. The process tree can be obtained by using the following command: *ps tree -s -p 2314*

```
cs1217@cs1217-devel:~$ echo $SHELL
/bin/bash
cs1217@cs1217-devel:~$ echo $$
2314
cs1217@cs1217-devel:~$ ps tree -s -p 2314
systemd(1)---systemd(1507)---gnome-terminal-(2240)---bash(2314)---pstree(2699)
cs1217@cs1217-devel:~$
```

-p flag is used to display the PIDs and -s shows parent processes of the specified process.

Question 6

The following system programs are simply **exec'ed** by the bash shell: **ls**, **ps**.

The system programs that are **implemented** by the bash code itself are: **cd**, **history**.

- cd**: `cd` is implemented by the bash code itself. Using the `type` command we can see that the program is builtin and the `strace` command for the same indicated that it is not exec'ed by the bash shell.
- ls**: `ls` is exec'ed by the bash shell. Using the `type` command with the `-a` flag (which deals with aliases) and `strace` clearly shows that `ls` is exec'ed by the bash code itself.
- history**: `history` is implemented by the bash code itself. Using the `type` command we can see that the program is builtin and the `strace` command for the same indicated that it is not exec'ed by the bash shell.
- ps**: `ps` is exec'ed by the bash shell. Using `type` command shows that `ps` is not a built-in command and `strace` clearly shows that `ps` command exec using the binary file located at the path of the command and then proceeds.

The following screenshots represents the output of `type` and `strace` command:

```
cs1217@cs1217-devel:~/Desktop/cs1217/cs1217-assignment-2-losethos$ type cd
cd is a shell builtin
cs1217@cs1217-devel:~/Desktop/cs1217/cs1217-assignment-2-losethos$ type ls
ls is aliased to 'ls --color=auto'
cs1217@cs1217-devel:~/Desktop/cs1217/cs1217-assignment-2-losethos$ type history
history is a shell builtin
cs1217@cs1217-devel:~/Desktop/cs1217/cs1217-assignment-2-losethos$ type ps
ps is /usr/bin/ps
cs1217@cs1217-devel:~/Desktop/cs1217/cs1217-assignment-2-losethos$ type ls
ls is aliased to 'ls --color=auto'
cs1217@cs1217-devel:~/Desktop/cs1217/cs1217-assignment-2-losethos$ type ls --color=auto
ls is aliased to 'ls --color=auto'
bash: type: --color=auto: not found
cs1217@cs1217-devel:~/Desktop/cs1217/cs1217-assignment-2-losethos$ type 'ls --color=auto'
bash: type: ls --color=auto: not found
cs1217@cs1217-devel:~/Desktop/cs1217/cs1217-assignment-2-losethos$ type ls
ls is aliased to 'ls --color=auto'
cs1217@cs1217-devel:~/Desktop/cs1217/cs1217-assignment-2-losethos$ type -a ls
ls is aliased to 'ls --color=auto'
ls is /usr/bin/ls
ls is /bin/ls
cs1217@cs1217-devel:~/Desktop/cs1217/cs1217-assignment-2-losethos$ man type
No manual entry for type
cs1217@cs1217-devel:~/Desktop/cs1217/cs1217-assignment-2-losethos$ type -p cd
cs1217@cs1217-devel:~/Desktop/cs1217/cs1217-assignment-2-losethos$ type -p ls
cs1217@cs1217-devel:~/Desktop/cs1217/cs1217-assignment-2-losethos$ type -p ps
/usr/bin/ps
cs1217@cs1217-devel:~/Desktop/cs1217/cs1217-assignment-2-losethos$ type -p history
cs1217@cs1217-devel:~/Desktop/cs1217/cs1217-assignment-2-losethos$ type -p -a ls
/usr/bin/ls
/bin/ls
cs1217@cs1217-devel:~/Desktop/cs1217/cs1217-assignment-2-losethos$ type -t ls
alias
cs1217@cs1217-devel:~/Desktop/cs1217/cs1217-assignment-2-losethos$ type -t -a ls
alias
file
file
cs1217@cs1217-devel:~/Desktop/cs1217/cs1217-assignment-2-losethos$ type -t ps
file
cs1217@cs1217-devel:~/Desktop/cs1217/cs1217-assignment-2-losethos$ type -t cd
builtin
cs1217@cs1217-devel:~/Desktop/cs1217/cs1217-assignment-2-losethos$ type -t history
builtin
cs1217@cs1217-devel:~/Desktop/cs1217/cs1217-assignment-2-losethos$
```

```

cs1217@cs1217-devel:~/Desktop/cs1217/cs1217-assignment-2-losethos$ strace -c ls
cpu  cpu.c  cpu-print  cpu-print.c  disk  disk1  disk1.c  disk.c  files  README.md
% time      seconds  usecs/call    calls    errors syscall
-----
67.21  0.009276      9276         1         write
8.53   0.001177         65        18         mmap
4.74   0.000654         93         7         openat
3.48   0.000480         53         9         close
3.22   0.000444         55         8         newfstatat
2.77   0.000383         63         6         mprotect
2.06   0.000285         57         5         read
1.28   0.000176         44         4         pread64
0.99   0.000136         68         2         getdents64
0.88   0.000122         61         2         statfs
0.87   0.000120         60         2         ioctl
0.78   0.000107         35         3         brk
0.52   0.000072         72         1         munmap
0.46   0.000063         31         2         access
0.39   0.000054         54         1         set_tid_address
0.38   0.000052         52         1         set_robust_list
0.38   0.000052         52         1         prlimit64
0.38   0.000052         52         1         rseq
0.37   0.000051         51         1         getrandom
0.33   0.000046         23         2         arch_prctl
0.00   0.000000         0          1         execve
-----
100.00  0.013802      176         78         5 total
cs1217@cs1217-devel:~/Desktop/cs1217/cs1217-assignment-2-losethos$ strace -c cd
strace: Can't stat 'cd': No such file or directory

```

```

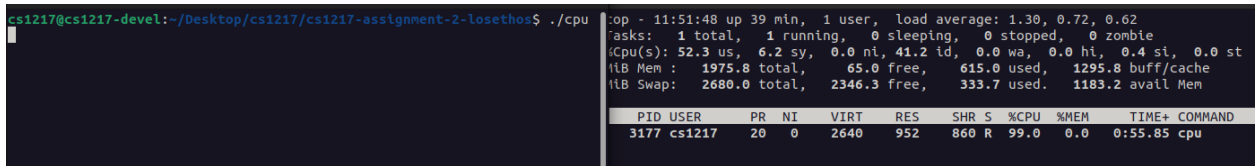
cs1217@cs1217-devel:~/Desktop/cs1217/cs1217-assignment-2-losethos$ strace -c ps
PID TTY      TIME CMD
2385 pts/0    00:00:00 bash
4503 pts/0    00:00:00 strace
4506 pts/0    00:00:00 ps
% time      seconds  usecs/call    calls    errors syscall
-----
44.37  0.004509         10        428         read
19.46  0.001978         4         434         openat
13.17  0.001338         3         430         close
11.05  0.001123         4         229         newfstatat
3.69   0.000375        187         2         getdents64
3.14   0.000319         6         50         mmap
1.98   0.000201        50         4         pread64
1.44   0.000146         9         15         mprotect
0.69   0.000070        17         4         write
0.46   0.000047         2         22         rt_sigaction
0.16   0.000016         16         1         munmap
0.08   0.000008         1         6         prctl
0.06   0.000006         3         2         ioctl
0.05   0.000005         1         3         brk
0.03   0.000003         1         2         lseek
0.03   0.000003         1         2         arch_prctl
0.03   0.000003         3         1         getrandom
0.02   0.000002         2         1         geteuid
0.02   0.000002         2         1         futex
0.02   0.000002         2         1         set_tid_address
0.02   0.000002         2         1         set_robust_list
0.02   0.000002         2         1         prlimit64
0.02   0.000002         2         1         rseq
0.00   0.000000         0         1         access
0.00   0.000000         0         1         execve
-----
100.00  0.010162         6        1643        13 total
cs1217@cs1217-devel:~/Desktop/cs1217/cs1217-assignment-2-losethos$ strace -c history
strace: Can't stat 'history': No such file or directory
cs1217@cs1217-devel:~/Desktop/cs1217/cs1217-assignment-2-losethos$

```

Question 7

- (a) **cpu**: The bottleneck resource for the `cpu.c` program is the **CPU**.

We ran the `top` command and executed the `cpu.c` file simultaneously. The `%CPU` field came out as around 100% for the `cpu` executable after running the command as visible in the following screenshot:



```
cs1217@cs1217-devel:~/Desktop/cs1217/cs1217-assignment-2-losestos$ ./cpu
top - 11:51:48 up 39 min, 1 user, load average: 1.30, 0.72, 0.62
tasks: 1 total, 1 running, 0 sleeping, 0 stopped, 0 zombie
%Cpu(s): 52.3 us, 6.2 sy, 0.0 ni, 41.2 id, 0.0 wa, 0.0 hi, 0.4 si, 0.0 st
MiB Mem : 1975.8 total, 65.0 free, 615.0 used, 1295.8 buff/cache
MiB Swap: 2680.0 total, 2346.3 free, 333.7 used, 1183.2 avail Mem

  PID USER      PR  NI  VIRT  RES  SHR  S  %CPU  %MEM    TIME+  COMMAND
 3177 cs1217    20   0   2640   952  860  R   99.0   0.0   0:55.85  cpu
```

After reading the code of `cpu.c`, we can see that it contains an infinite while loop which in turn contains a for loop that calculates factorials. This process is a very resource heavy task using ALU unit of CPU. It keeps on executing this intensive for loop forever without any breaks in between which leads it to maximise CPU usage.

- (b) **cpu-print**: The bottleneck resource for the `cpu-print.c` program is the **I/O**.

We ran the `vmstat` command multiple times at different intervals after executing `cpu-print`, and we can observe that the number in `bi` column (blocks read from the disk) is reducing, as visible in the following screenshot. Also, we can observe that when we run `top` command and check the details about various processes, `cpu-print` is switching state between `R` and `S`, and `gnome-terminal` along with `cpu-print` is consuming most of the `%CPU` resource, thus CPU acts as a secondary bottleneck

After reading the code of `cpu-print.c`, we can see that the program is making a system call to get the time of the day and to print the same. Since, it is printing continuously on the screen, the I/O is blocked by the input it is getting from the system calls, making I/O as the bottleneck.


```

cs1217@cs1217-devel: ~/Desktop/cs1217/cs1217-assignment-2-losethos

cs1217@cs1217-devel:~$ vmstat
procs -----memory----- --swap-- ----io---- -system-- -----cpu-----
r b swpd free buff cache si so bi bo in cs us sy id wa st
1 0 0 454588 40084 750284 0 0 1536 58 373 2917 4 10 85 1 0

cs1217@cs1217-devel:~$ vmstat
procs -----memory----- --swap-- ----io---- -system-- -----cpu-----
r b swpd free buff cache si so bi bo in cs us sy id wa st
1 0 0 454588 40092 750548 0 0 1507 58 425 5346 5 11 84 1 0

cs1217@cs1217-devel:~$ vmstat
procs -----memory----- --swap-- ----io---- -system-- -----cpu-----
r b swpd free buff cache si so bi bo in cs us sy id wa st
1 0 0 454588 40092 750396 0 0 1485 58 464 7755 5 11 83 1 0

cs1217@cs1217-devel:~$ vmstat
procs -----memory----- --swap-- ----io---- -system-- -----cpu-----
r b swpd free buff cache si so bi bo in cs us sy id wa st
2 0 0 454588 40100 750308 0 0 1473 58 480 8906 5 12 83 1 0

cs1217@cs1217-devel:~$ vmstat
procs -----memory----- --swap-- ----io---- -system-- -----cpu-----
r b swpd free buff cache si so bi bo in cs us sy id wa st
4 0 0 454600 40164 750384 0 0 1326 58 707 22311 5 15 79 1 0

cs1217@cs1217-devel:~$ vmstat
procs -----memory----- --swap-- ----io---- -system-- -----cpu-----
r b swpd free buff cache si so bi bo in cs us sy id wa st
3 0 0 452736 40484 758064 0 0 1053 57 1068 45027 7 23 69 1 0

cs1217@cs1217-devel:~$ vmstat
procs -----memory----- --swap-- ----io---- -system-- -----cpu-----
r b swpd free buff cache si so bi bo in cs us sy id wa st
3 0 0 452488 40692 758364 0 0 806 58 1483 21431 8 30 61 0 0

cs1217@cs1217-devel:~$ vmstat
procs -----memory----- --swap-- ----io---- -system-- -----cpu-----
r b swpd free buff cache si so bi bo in cs us sy id wa st
1 0 0 452488 40768 758076 0 0 746 57 1574 29801 8 32 59 0 0

cs1217@cs1217-devel:~$ vmstat
procs -----memory----- --swap-- ----io---- -system-- -----cpu-----
r b swpd free buff cache si so bi bo in cs us sy id wa st
1 0 0 452488 40768 758280 0 0 744 57 1578 30184 9 32 59 0 0

cs1217@cs1217-devel:~$ vmstat
procs -----memory----- --swap-- ----io---- -system-- -----cpu-----
r b swpd free buff cache si so bi bo in cs us sy id wa st
2 0 0 452488 40932 758224 0 0 698 58 1655 36895 9 33 58 0 0

cs1217@cs1217-devel:~$ vmstat
procs -----memory----- --swap-- ----io---- -system-- -----cpu-----
r b swpd free buff cache si so bi bo in cs us sy id wa st
1 0 0 452488 40940 758308 0 0 696 58 1658 37008 9 33 58 0 0

```

(c) **disk**: The bottleneck resource for the disk.c program is the **Disk**.

We ran the atop command and executed the disk.c file simultaneously. We observe that the program keeps switching state (S) between R(running) and D(uninterrupted-sleep, mainly caused due to disk full). Also, the DSK field came out as around 115% for the disk executable after running the command as visible in the following screenshot:

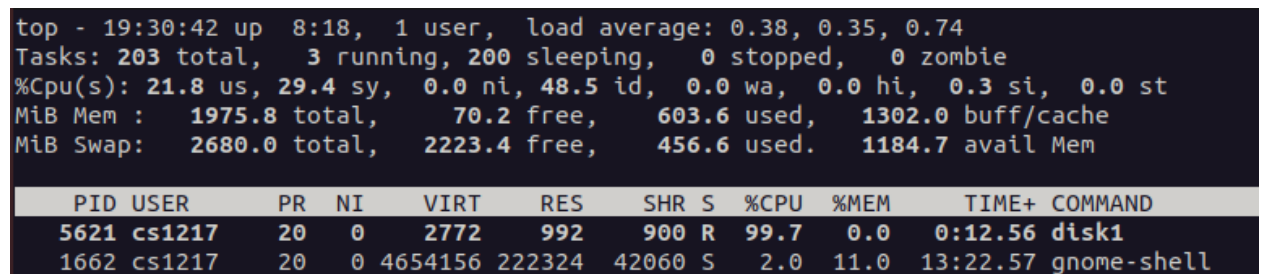
cs1217@cs1217-devel: ~/Desktop/cs1217/cs1217-assignment-2-losethos															cs1217@cs1217-devel: ~/Desktop/cs1217/cs1217-assignment-2-losethos														
TOP - cs1217-devel															2023/02/14 19:05:41														
RC	sys	7.78s	user	1.03s						#proc	208	#lrun	1	#slp	425	#tslp	49	#zombie	0	clones	0	ldest	0						
PU	sys	61%	user	9%		lrq	41%			idle	70%	wait	19%					steal	0%	guest	0%	curf	2.61GHz						
PU	sys	41%	user	4%		lrq	34%			idle	14%	cpu001	7%					steal	0%	guest	0%	curf	2.61GHz						
PU	sys	22%	user	5%		lrq	9%			idle	52%	cpu000	12%					steal	0%	guest	0%	curf	2.61GHz						
PL	avg1	2.06			avg5	1.70	avg15	1.18						cs	19363		lntr	21310			nuncpu	2							
EN	tot	1.96	free	74.9M	cache	1.2G	dirty	0.0M	buff	12.6M	slab	141.4M	slrec	61.4M	shmem	16.6M	shrss	0.0M	shswp	0.0M	vmcon	3.8G	vmlln	3.6G					
UP	tot	2.6G			free	2.2G	swcac	54.2M									swin	0	swout	4242	oomkill	0							
SI	scan	664382	steal	624630	stall	0	compact	0		numantg	0	migrate	14e3	cs	8/13/14	ns	3/1/0	mf	3/1/0	ls	17/6/2	tf	15/5/1						
	cpusone	9%	memone	4%	memfull	4%				iosone	22%	lofull	20%	cs	8/13/14	ns	3/1/0	mf	3/1/0	ls	17/6/2	tf	15/5/1						
	sda	busy	115%		read	12370	write	449		discrd	0	KlB/r	100	KlB/w	37			MBw/s	121.5	MBw/s	1.7	avg	1.40	avto	0.76 ns				
PID	SYSCPU	USRCPU	RDELAY	VGROW	RGROW	RUID	EUID	ST	EXC	THR	S	CPUNR	CPU	CMD	1/4														
5581	5.60s	0.85s	0.61s	0B	0B	cs1217	cs1217	--	-	1	D	1	66%	disk															
1602	0.36s	0.79s	0.88s	7.0M	-17.9M	cs1217	cs1217	--	-	14	S	0	13%	gnome-shell															
92	1.06s	0.88s	0.07s	0B	0B	root	root	--	-	1	S	0	12%	kswapd0															

After reading the code of disk.c, we can see that the program is randomly opening any file from

the files folder, reading a block of data in 1024 bytes chunks by putting it into buffer and then discards the buffer. Since, there is an infinite loop which keeps opening random files, the disk space runs out quickly, thus becoming a bottleneck for the process.

- (d) **disk1**: The bottleneck resource for the disk1.c program is the **CPU**.

We ran the top command and executed the disk1.c file simultaneously. The %CPU field came out as around 100% for the disk1 executable after running the command as visible in the following screenshot:



```
top - 19:30:42 up 8:18, 1 user, load average: 0.38, 0.35, 0.74
Tasks: 203 total, 3 running, 200 sleeping, 0 stopped, 0 zombie
%Cpu(s): 21.8 us, 29.4 sy, 0.0 ni, 48.5 id, 0.0 wa, 0.0 hi, 0.3 si, 0.0 st
MiB Mem : 1975.8 total, 70.2 free, 603.6 used, 1302.0 buff/cache
MiB Swap: 2680.0 total, 2223.4 free, 456.6 used. 1184.7 avail Mem
```

PID	USER	PR	NI	VIRT	RES	SHR	S	%CPU	%MEM	TIME+	COMMAND
5621	cs1217	20	0	2772	992	900	R	99.7	0.0	0:12.56	disk1
1662	cs1217	20	0	4654156	222324	42060	S	2.0	11.0	13:22.57	gnome-shell

After reading the code of disk1.c, we can see that it contains an infinite while loop which opens and reads foo0.txt from the files subfolder. Like disk.c, it reads the data in the file in blocks of 1024 bytes. But it doesn't delete/remove the data it reads, which fills up the CPU quickly, hence making it a bottleneck for this program.