

# CS1217 - Spring 2023 - Lab 1

Bhumika Mittal, Saptarishi Dhanuka

## Part 1: PC Bootstrap

```
bhumika@bhumika:~/Desktop/cs1217/cs1217-lab-1-losethos/jos$ make qemu-nox
***
*** Use Ctrl-a x to exit qemu
***
qemu-system-i386 -nographic -drive file=obj/kern/kernel.img,index=0,media=disk,format=raw -serial mon
:stdio -gdb tcp::26000 -D qemu.log
6828 decimal is XXX octal!
entering test_backtrace 5
entering test_backtrace 4
entering test_backtrace 3
entering test_backtrace 2
entering test_backtrace 1
entering test_backtrace 0
leaving test_backtrace 0
leaving test_backtrace 1
leaving test_backtrace 2
leaving test_backtrace 3
leaving test_backtrace 4
leaving test_backtrace 5
Welcome to the JOS kernel monitor!
Type 'help' for a list of commands.
K> help
help - Display this list of commands
kerninfo - Display information about the kernel
K> kerninfo
Special kernel symbols:
_start                0010000c (phys)
_entry f010000c (virt) 0010000c (phys)
etext f01019e1 (virt) 001019e1 (phys)
edata f0112060 (virt) 00112060 (phys)
end f01126c0 (virt) 001126c0 (phys)
Kernel executable memory footprint: 74KB
K> 
```

Make qemu: help and kerninfo

## The ROM BIOS

```
bhumika@bhumika: ~/Desktop/cs1217/cs1217-lab-1-loseth... x bhumika@bhumika: ~/Desktop/cs1217/cs1217-lab-1-loseth... x
bhumika@bhumika:~/Desktop/cs1217/cs1217-lab-1-loseth...$ cd jos
bhumika@bhumika:~/Desktop/cs1217/cs1217-lab-1-loseth.../jos$ make gdb
gdb -n -x .gdbinit
GNU gdb (Ubuntu 12.1-0ubuntu1~22.04) 12.1
Copyright (C) 2022 Free Software Foundation, Inc.
License GPLv3+: GNU GPL version 3 or later <http://gnu.org/licenses/gpl.html>
This is free software: you are free to change and redistribute it.
There is NO WARRANTY, to the extent permitted by law.
Type "show copying" and "show warranty" for details.
This GDB was configured as "x86_64-linux-gnu".
Type "show configuration" for configuration details.
For bug reporting instructions, please see:
<https://www.gnu.org/software/gdb/bugs/>.
Find the GDB manual and other documentation resources online at:
<http://www.gnu.org/software/gdb/documentation/>.

For help, type "help".
Type "apropos word" to search for commands related to "word".
+ target remote localhost:26000
warning: No executable has been specified and target does not support
determining executable automatically. Try using the "file" command.
warning: A handler for the OS ABI "GNU/Linux" is not built into this configuration
of GDB. Attempting to continue with the default i8086 settings.

The target architecture is set to "i8086".
[f000:ffff] 0xfffff0: jmp 0xf000,0xe05b
0x0000ffff in ?? ()
+ symbol-file kernel
(gdb) □
```

What the BIOS does with its first few instructions is initialize various registers like dx (data register), ss (stack segment register), esp, edx, then jumps to an earlier location in the BIOS. cli clears the interrupt flag, thus ensuring no interrupts occur and cld clears the direction flag, thus ensuring that string operations increment forward. Then out and in instructions are used for talking with hardware devices through the use of ports and registers like al. Ports about the CMOS RAM Real Time Clock, Data Port and system devices are used as per their port numbers. Then it loads the interrupt descriptor table and the global descriptor table registers and sets the first bit of control register cr0 to 1, which enables protected mode. It then jumps to a memory location and switches to 32 bit mode.

## Part 2: The Boot Loader

**Exercise 3.** *At what point does the processor start executing 32-bit code? What exactly causes the switch from 16- to 32-bit mode?*

In the xv6 operating system, the switch from 16-bit to 32-bit mode happens during the boot process. The boot loader switches the processor from real mode to 32-bit protected mode, because it is only in this mode that software can access all the memory above 1MB in the processor's physical address space.

When the boot loader loads the kernel image into memory, it sets the processor's mode to 32-bit protected mode before jumping to the kernel's entry point.

The exact point where the processor start executing 32-bit mode is the line 60 in the boot.S file

```
movw    $PROT_MODE_DSEG, %ax
```

In real mode, the PE bit in the cr0 control register is set to 0, indicating that the processor is operating in 16-bit real mode. The following instruction switches the PE bit to 1 in the cr0, thus enabling the 32-bit protected mode.

```
mov    %eax,%cr0
```

The following line (line 55) in jos/boot/boot.S file, switches processor into 32-bit mode.

```
ljmp    $PROT_MODE_CSEG, $protcseg
```

*What is the last instruction of the boot loader executed, and what is the first instruction of the kernel it just loaded?*

After the for loop which calls readseg(), the last instruction of the bootloader in main.c is

```
((void (*)(void)) (ELFHDR->e_entry))();
```

In boot.asm, the last instruction is right below the above line as:

```
7d71: ff 15 18 00 01 00    call    *0x10018
```

Since the last instruction of the bootloader is calling \*0x10018, we must see what is at this location. We first set a breakpoint at 7d71 and then single step.

```
Breakpoint 1, 0x00007c00 in ?? ()
(gdb) br * 0x7d71
Breakpoint 2 at 0x7d71
(gdb) c
Continuing.
The target architecture is set to "i386".
=> 0x7d71:    call    *0x10018

Breakpoint 2, 0x00007d71 in ?? ()
(gdb) si
=> 0x10000c:  movw    $0x1234,0x472
0x0010000c in ?? ()
```

Hence the first instruction of the kernel it just loaded is:

```
(gdb) si
=> 0x10000c:  movw    $0x1234,0x472
```

(gdb) x/i \*0x10018 also gives:

```
0x10000c:    movw    $0x1234,0x472
```

*Where is the first instruction of the kernel?*

As shown above the first instruction of the kernel is at **0x10000c**

*How does the boot loader decide how many sectors it must read in order to fetch the entire kernel from disk? Where does it find this information?*

The necessary information for the boot loader is given in the Executable and Linkable Format (ELF) header. In main.c, we define ELFHDR as a pointer to struct ELF with value 0x1000. Then, after calling readseg once, we define eph to be a pointer to a struct Proghdr which points to the end of the program header table since we used e\_phnum to define it (e\_phnum is number of elements in program header table which describes information that the bootloader needs to prepare for the kernel to execute).

Hence, it decides how many sectors based on the value of the phnum field in ELFHDR and simply loops through them all.

## Loading the Kernel

**Exercise 5.** *Identify the first instruction that would "break" or otherwise do the wrong thing if you were to get the boot loader's link address wrong.*

We changed the Makefrag to read, instead of 7C00,

```
-Ttext 0x7D00
```

Then after make clean and make, we debugged with gdb by setting a breakpoint at 0x7c00 and continuing execution.

By single stepping, the first instruction to be different was :

**0x7c1e: lgdtw 0x7d64** in the 7D00 case, whereas it was **0x7c1e: lgdtw 0x7c64** in the proper 7C00 case. We can see that the argument for lgdtw has changed from **0x7c64** to **0x7d64**, which is wrong. However, we can still continue execution from here.

The following instruction breaks the whole thing completely:

**0x7c2d: ljmp \$0x8,\$0x7d32**

It give the following error:

```
ld: warning: section '.bss' type changed to PROGBITS
+ as boot/boot.S
+ cc -Os boot/main.c
+ ld boot/boot
boot block is 396 bytes (max 510)
+ mk obj/kern/kernel.img
xv6@xv6host:~/lab1/cs1217-lab-1-losethos/jos$ make qemu-nox-gdb
sed "s/localhost:1234/localhost:26000/" < .gdbinit.tpl > .gdbinit
***
*** Now run 'make gdb'.
***
qemu-system-i386 -nographic -drive file=obj/kern/kernel.img,index=0,media=disk,format=
raw -serial mon:stdio -gdb tcp::26000 -D qemu.log -S
EAX=00000011 EBX=00000000 ECX=00000000 EDX=00000080
ESI=00000000 EDI=00000000 EBP=00000000 ESP=00006f20
EIP=00007c2d EFL=00000006 [-----P-] CPL=0 II=0 A20=1 SMM=0 HLT=0
ES =0000 00000000 0000ffff 00009300 DPL=0 DS16 [-WA]
CS =0000 00000000 0000ffff 00009b00 DPL=0 CS16 [-RA]
SS =0000 00000000 0000ffff 00009300 DPL=0 DS16 [-WA]
DS =0000 00000000 0000ffff 00009300 DPL=0 DS16 [-WA]
FS =0000 00000000 0000ffff 00009300 DPL=0 DS16 [-WA]
GS =0000 00000000 0000ffff 00009300 DPL=0 DS16 [-WA]
LDT=0000 00000000 0000ffff 00008200 DPL=0 LDT
TR =0000 00000000 0000ffff 00008b00 DPL=0 TSS32-busy
GDT= 006ee8ec 000073ff
IDT= 00000000 000003ff
CR0=00000011 CR2=00000000 CR3=00000000 CR4=00000000
DR0=00000000 DR1=00000000 DR2=00000000 DR3=00000000
DR6=ffff0fff DR7=00000400
EFER=0000000000000000
Triple fault. Halting for inspection via QEMU monitor.

The target architecture is set to "i8086".
[f000:fff0] 0xffff0: ljmp $0xf000,$0xe05b
0x0000fff0 in ?? ()
+ symbol-file kernel
(gdb) br * 0x7c00
Breakpoint 1 at 0x7c00
(gdb) c
Continuing.
[ 0:7c00] => 0x7c00: cli
Breakpoint 1, 0x00007c00 in ?? ()
(gdb) c
Continuing.
Program received signal SIGTRAP, Trace/breakpoint trap.
[ 0:7c2d] => 0x7c2d: ljmp $0x8,$0x7d32
0x00007c2d in ?? ()
(gdb) █
```

The lgdtw instruction gets a wrong argument hence a wrong value is loaded into the global descriptor table register

The ljmp instruction which switches from 16 bits to 32 bits breaks everything if the wrong link address is given in boot/Makefrag

Also note that, after the above instruction, the bootloader goes into the infinite spin loop (which only gets executed when the bootmain returns

OUTPUT	DEBUG CONSOLE	TERMINAL	GITLENS
(gdb) si	[ 0:7c26] => 0x7c26: or \$0x1,%eax		
0x00007c26 in ?? ()			
(gdb) si	[ 0:7c2a] => 0x7c2a: mov %eax,%cr0		
0x00007c2a in ?? ()			
(gdb) si	[ 0:7c2d] => 0x7c2d: ljmp \$0x8,\$0x7d32		
0x00007c2d in ?? ()			
(gdb) si	[ 0:7c2d] => 0x7c2d: ljmp \$0x8,\$0x7d32		
0x00007c2d in ?? ()			
(gdb) si	[ 0:7c2d] => 0x7c2d: ljmp \$0x8,\$0x7d32		
0x00007c2d in ?? ()			
(gdb) si	[ 0:7c2d] => 0x7c2d: ljmp \$0x8,\$0x7d32		
0x00007c2d in ?? ()			
(gdb) si	[ 0:7c2d] => 0x7c2d: ljmp \$0x8,\$0x7d32		
0x00007c2d in ?? ()			
(gdb) █			

**Exercise 6.** Examine the 8 words of memory at `0x00100000` at the point the BIOS enters the boot loader, and then again at the point the boot loader enters the kernel. Why are they different? What is there at the second breakpoint?

First we set a breakpoint at `0x7c00` (where the BIOS enters the boot loader) and continue to it and single step. We see that 8 words of memory around the given memory location are all 0

```
(gdb) br * 0x7c00
Breakpoint 1 at 0x7c00
(gdb) c
Continuing.
[ 0:7c00] => 0x7c00: cli

Breakpoint 1, 0x00007c00 in ?? ()
(gdb) x/8x 0x00100000
0x100000: 0x00000000 0x00000000 0x00000000 0x00000000
0x100010: 0x00000000 0x00000000 0x00000000 0x00000000
(gdb) si
[ 0:7c01] => 0x7c01: cld
0x00007c01 in ?? ()
(gdb) x/8x 0x00100000
0x100000: 0x00000000 0x00000000 0x00000000 0x00000000
0x100010: 0x00000000 0x00000000 0x00000000 0x00000000
```

Then we set a breakpoint at `0x7d71`, which is where the boot loader enters the kernel. Examining 8 words of memory shows change.

```
Breakpoint 2 at 0x7d71
(gdb) c
Continuing.
The target architecture is set to "i386".
=> 0x7d71: call *0x10018

Breakpoint 2, 0x00007d71 in ?? ()
(gdb) x/8x 0x00100000
0x100000: 0x1badb002 0x00000000 0xe4524ffe 0x7205c766
0x100010: 0x34000004 0x1000b812 0x220f0011 0xc0200fd8
(gdb) si
=> 0x10000c: movw $0x1234,0x472
0x0010000c in ?? ()
(gdb) x/8i 0x00100000
0x100000: add 0x1bad(%eax),%dh
0x100006: add %al,(%eax)
0x100008: decb 0x52(%edi)
0x10000b: in $0x66,%al
0x10000d: movl $0xb81234,0x472
0x100017: adc %dl,(%ecx)
0x100019: add %cl,(%edi)
0x10001b: and %al,%bl
(gdb) x/8x 0x00100000
0x100000: 0x1badb002 0x00000000 0xe4524ffe 0x7205c766
0x100010: 0x34000004 0x1000b812 0x220f0011 0xc0200fd8
```

At the second breakpoint, the boot loader has loaded the kernel into memory starting from the kernel's load address `0x00100000`, and hence the 8 words we see are the first few words of memory of the kernel program. Meanwhile, at the first breakpoint, the BIOS has just entered the boot loader and hence there is nothing at address `0x00100000`, which is why they are different.

## Part 3: The Kernel

### Using virtual memory to work around position dependence

**Exercise 7.** What is the first instruction after the new mapping is established that would fail to work properly if the mapping weren't in place?

The following screenshot shows that after `mov %eax, %cr0`, a mapping has been established between `0x00100000` and `0xf0100000`, along with the other memory location ranges mentioned in the question. The JOS kernel has set up a mapping between the virtual address `0xf0100000` and the physical address `0x00100000`, so any memory accesses to `0xf0100000` will be translated to `0x00100000`. Here, memory references are getting translated by the virtual memory hardware to physical addresses.

```
(gdb) si
=> 0x100020:  or      $0x80010001,%eax
0x00100020 in ?? ()
(gdb) x/8x 0x00100000
0x100000:  0x1badb002      0x00000000      0xe4524ffe      0x7205c766
0x100010:  0x34000004      0x1000b812      0x220f0011      0xc0200fd8
(gdb) x/5x 0x00100000
0x100000:  0x1badb002      0x00000000      0xe4524ffe      0x7205c766
0x100010:  0x34000004
(gdb) x/4x 0x00100000
0x100000:  0x1badb002      0x00000000      0xe4524ffe      0x7205c766
(gdb) x/4x 0xf0100000
0xf0100000 <_start-268435468>: 0x00000000      0x00000000      0x00000000      0x00000000
(gdb) si
=> 0x100025:  mov      %eax,%cr0
0x00100025 in ?? ()
(gdb) x/4x 0x00100000
0x100000:  0x1badb002      0x00000000      0xe4524ffe      0x7205c766
(gdb) x/4x 0xf0100000
0xf0100000 <_start-268435468>: 0x00000000      0x00000000      0x00000000      0x00000000
(gdb) si
=> 0x100028:  mov      $0xf010002f,%eax
0x00100028 in ?? ()
(gdb) x/4x 0x00100000
0x100000:  0x1badb002      0x00000000      0xe4524ffe      0x7205c766
(gdb) x/4x 0xf0100000
0xf0100000 <_start-268435468>: 0x1badb002      0x00000000      0xe4524ffe      0x7205c766
```

The first instruction that wouldn't work as it is supposed to is:

```
jmp *%eax
```

After commenting out `movl %eax, %cr0`, running *make clean* then *make* and using *gdb*, we get the following results and error:

```
PROBLEMS OUTPUT DEBUG CONSOLE JUPYTER PORTS TERMINAL
+ cc kern/kdebug.c
+ cc lib/printfmt.c
+ cc lib/readline.c
+ cc lib/string.c
+ ld obj/kern/kernel
ld: warning: section '.bss' type changed to PROGBITS
+ as boot/boot.S
+ cc -Os boot/main.c
+ ld boot/boot
boot block is 396 bytes (max 510)
+ mk obj/kern/kernel.img
xv6@xv6host:~/lab1/cs1217-lab-1-losesthos/jos$ make qemu-nox-gdb
sed "s/localhost:1234/localhost:26000/" < .gdbinit.tmpl > .gdbinit
***
*** Now run 'make gdb'.
***
qemu-system-i386 -nographic -drive file=obj/kern/kernel.img,index=0,media=disk,format=
raw -serial mon:stdio -gdb tcp::26000 -D qemu.log -S
qemu: fatal: Trying to execute code outside RAM or ROM at 0xf010002c

EAX=f010002c EBX=00010094 ECX=00000000 EDX=ffffff40
ESI=00010094 EDI=00000000 EBP=00007bf8 ESP=00007bec
EIP=f010002c EFL=00000086 [---P---] CPL=0 II=0 A20=1 SMM=0 HLT=0
ES =0010 00000000 ffffffff 00cf9300 DPL=0 DS [-WA]
CS =0008 00000000 ffffffff 00cf9a00 DPL=0 CS32 [-R-]
SS =0010 00000000 ffffffff 00cf9300 DPL=0 DS [-WA]
DS =0010 00000000 ffffffff 00cf9300 DPL=0 DS [-WA]
FS =0010 00000000 ffffffff 00cf9300 DPL=0 DS [-WA]
GS =0010 00000000 ffffffff 00cf9300 DPL=0 DS [-WA]
LDT=0000 00000000 0000ffff 00000200 DPL=0 LDT
TR =0000 00000000 0000ffff 00000b00 DPL=0 TSS32-busy
GDT= 00007c4c 00000017
IDT= 00000000 000003ff
CR0=00000011 CR2=00000000 CR3=00111000 CR4=00000000
DR0=00000000 DR1=00000000 DR2=00000000 DR3=00000000
DR6=ffff0f10 DR7=00000000
CCS=00000084 CCD=80010011 CCO=EFLAGS
EFER=0000000000000000
FCW=037f FSW=0000 [ST=0] FTW=0 MXCSR=00001f80
FPR0=0000000000000000 0000 FPR1=0000000000000000 0000
FPR2=0000000000000000 0000 FPR3=0000000000000000 0000
FPR4=0000000000000000 0000 FPR5=0000000000000000 0000
FPR6=0000000000000000 0000 FPR7=0000000000000000 0000
XMM0=00000000000000000000000000000000 XMM01=00000000000000000000000000000000
XMM02=00000000000000000000000000000000 XMM03=00000000000000000000000000000000
XMM04=00000000000000000000000000000000 XMM05=00000000000000000000000000000000
XMM06=00000000000000000000000000000000 XMM07=00000000000000000000000000000000
make: *** [GNUmakefile:174: qemu-nox-gdb] Aborted (core dumped)
xv6@xv6host:~/lab1/cs1217-lab-1-losesthos/jos$

+ symbol-file kernel
(gdb) br * 0x7d71
Breakpoint 1 at 0x7d71
(gdb) c
Continuing.
The target architecture is set to "i386".
=> 0x7d71: call *0x10018

Breakpoint 1, 0x00007d71 in ?? ()
(gdb) si
=> 0x10000c: movw $0x1234,0x472
0x0010000c in ?? ()
(gdb) si
=> 0x100015: mov $0x111000,%eax
0x00100015 in ?? ()
(gdb) si
=> 0x10001a: mov %eax,%cr3
0x0010001a in ?? ()
(gdb) info reg eax
eax 0x111000 1118208
(gdb) si
=> 0x10001d: mov %cr0,%eax
0x0010001d in ?? ()
(gdb) info reg eax
eax 0x111000 1118208
(gdb) si
=> 0x100020: or $0x80010001,%eax
0x00100020 in ?? ()
(gdb) info reg eax
eax 0x11 17
(gdb) si
=> 0x100025: mov $0xf010002c,%eax
0x00100025 in ?? ()
(gdb) si
=> 0x10002a: jmp *%eax
0x0010002a in ?? ()
(gdb) info reg eax
eax 0xf010002c -267386836
(gdb) x/4x 0xf010002c
0xf010002c <relocated>: 0x00000000 0x00000000 0x00000000 0x00000000
(gdb) x/4x 0x0010002c
0x10002c: 0x000000bd 0xf000bc00 0x68e8f010 0xeb000000
(gdb) si
=> 0xf010002c <relocated>: add %al,(%eax)
relocated () at kern/entry.S:74
74 movl $0x0,%ebp # nuke frame pointer
(gdb) si
Remote connection closed
(gdb)
```

This is because we first move the address of relocated to eax, which is 0xf010002c, which we can see does not have anything in it since mapping has not been established. Then we are jumping to that address, which is incorrect hence leading to an error.

Without the virtual memory mapping in place, this instruction would set the stack pointer to an incorrect physical address, causing memory access errors and ultimately leading to the segmentation fault.



## Formatted Printing to the Console

**Exercise 8.** A small fragment of code has been omitted - the code necessary to print octal numbers using patterns of the form "%o". Find and fill in this code fragment.

Added the following code necessary to print octal numbers using patterns of the form "%o".

```
num = getuint(&ap, lflag);
if (altflag && num != 0){
    putchar('0', putdat);
}
base = 8;
goto number;
```

Explain the interface between `printf.c` and `console.c`. Specifically, what function does `console.c` export? How is this function used by `printf.c`?

`printf.c` is responsible for the implementation of the `cprintf()` function which prints output to the console. The function that `console.c` exports that is used by `printf.c` is the **`cputchar()`** function, which in turn just calls the `cons_putc()` function, which is also defined in `console.c`

`printf.c` uses the `cputchar()` function in the following way: when the user calls the `cprintf()` function, it calls another function `vcprintf()` which formats the string using `vprintfmt()`, and outputs it to the console using the `putch()` function, which ultimately calls `cputchar()`.

Therefore, `console.c` manages the printing of the characters to the console whereas `printf.c` formats the output and passes it to the `console.c` to display the output. `cputchar()` and (indirectly) `consputc()` is used by `printf.c` to **output each character to the console**.

Explain the code from `console.c`:

This code is a very simple implementation of scrolling the console display. The main purpose of this code is if the cursor is going off the screen, scroll up.

If the current position of the cursor (`crt_pos`) is greater than or equal to the total size of the console buffer (`CRT_SIZE`), then we need to scroll up because the display is at the bottom of the screen. To do this, the `memmove()` function (it copy `n` bytes from `src` to `dst`) takes `crt_buf + CRT_COLS` as source, and moves  $(\text{CRT\_SIZE} - \text{CRT\_COLS}) * \text{sizeof}(\text{uint16\_t})$  number of bytes to `crt_buf`. Here,  $(\text{CRT\_SIZE} - \text{CRT\_COLS}) * \text{sizeof}(\text{uint16\_t})$  is the size of the remaining console buffer. Basically, this is moving contents of the console buffer up by one row.

Then the for loop effectively clears the last row of the console buffer and sets its attributes to `0x0700` (white text on black background). The loop iterates through the last row, setting each character to a space.

*Trace the execution of the following code step-by-step:*

```
int x = 1, y = 3, z = 4;
cprintf("x %d, y %x, z %d\n", x, y, z);
```

We can put this code in monitor.c just after

```
cprintf("Type 'help' for a list of commands.\n");
```

and use gdb to step into the function to help us to see the step by step execution of the code. We can also consult kernel.asm

Output of code: "x 1, y 3, z 4"

Step by step execution of the code:

- a. Declare and initialize x=1, y=3, and z=4
- b. Call cprintf function with the arguments: format string "x %d, y %x, z %d\n" and the variables x,y,z
- c. Declare a va\_list (a type defined in stdarg.h that is used to iterate through a variable number of arguments passed to a function) ap and an integer cnt.
- d. Call vprintf with arguments as the format string and the argument list.
- e. vprintf will iterate over the format string and calls putch function for each character. When %d is encountered, then it calls printnum to print the number in decimal format. Similarly for %x and %o it prints the number in hexadecimal and octal format. For \n, it calls putch
- f. printnum function calls putch for each character in the string (the integer argument is converted to string)
- g. putch calls cons\_putc from console.c and prints each character to the console. It also checks if the character is a newline. If it is a newline character, it updates the crt\_pos by one (if it is at the end of the display, it scrolls up).
- h. cons\_putc writes to the console buffer at the current crt\_pos and keeps updating it.
- i. vprintf then returns and transfers the control back to the caller, cprintf, which also then returns.

NOTE - va\_start initializes the va\_list to the first argument after fmt and va\_end cleans up the va\_list.

**fmt** points to the format string "x %d, y %x, z %d\n" and **ap** points to argument x since va\_list is initialized as the first argument

We used gdb for this part along with logically tracing the code and finding where the functions would execute and what their arguments would be:

- `vcprintf(fmt, ap)`  
`fmt` points to `"x %d, y %x, z %d\n"` and `ap` points to the value in `x`. Gdb gave me the values as `vcprintf (fmt=0xf0101d0e "x %d, y %x, z %d\n", ap=0xf010ef54 "\001")` since the value in `x` is 1
- `cons_putc()` with argument 120, which is `x` in ASCII
- `cons_putc()` with argument 32 // ASCII for space
- Then call to `va_arg`. Before the call, `ap` points to `"\001"`, but the next call will set `ap` to the next number to be printed i.e. `"\003"`
- `cons_putc(49)` // ascii for 1
- Then `cons_putc()` is called again and again with arguments as the ascii values for comma, space, y, space.
- `va_arg()`  
Before: points to `"\003"`  
After: points to `"\004"`
- `cons_putc()` called multiple times with arguments as ascii values for 3, comma, space, z, space
- `va_arg()`  
Before: points to `"\004"`  
After: Garbage value
- `cons_putc()` called twice with arguments as ascii values for 4 and newline.

For `va_arg()` answers we checked the values of the memory locations with gdb starting from the initial value of `ap = 0xf010ef54` till `0xf010ef60` and they checked out with the logical analysis

```

0xf010ef54: 0x01 0x00 0x00 0x00
(gdb) x/4d 0xf010ef54
0xf010ef54: 1 0 0 0
(gdb) x/4d 0xf010ef58
0xf010ef58: 3 0 0 0
(gdb) x/4d 0xf010ef5c
0xf010ef5c: 4 0 0 0
(gdb) x/4d 0xf010ef60

```

NOTE – We then removed the lines we added to `monitor.c` and didn't commit it

*What is the output?*

Output – **He110 World**

57616 in decimal is E110 in hex and 0x00646c72 is stored as 72 6c 64 00 which is the ASCII for `"rld\0"`.

If the x86 were instead big-endian, we would need to reverse the order of bytes in `i` such that it still contains the ASCII for `"rld\0"`. Converting 0x00646c72 to big-endian, we get 0x726c6400. In order to yield the same output, set `i = 0x726c6400`. There is **no need to change 57616** to a different value because this value is passed as an argument to `cprintf` in a way such that it prints this unsigned hexadecimal integer as ASCII. It is not affected by the byte order of x86.

*In the following code, what is going to be printed after 'y='? (note: the answer is not a specific value.) Why does this happen?*

Output: x=3 y=734328248

The format string contains two integer placeholders but only one int argument is provided. Observe that the first placeholder correctly takes the value 3 but the second one takes a **garbage value**. When the code is executed, it sometimes prints a garbage value or can even crash the program sometimes.

*How would you have to change `cprintf` or its interface so that it would still be possible to pass it a variable number of arguments?*

One way to do this is to pass the number of arguments to `cprintf` as the first parameter, before the format string. We can then iterate through the arguments using a loop inside `cprintf` and the `va_arg` macro, which will allow us to retrieve each argument from the variable argument list. By this method, we would be able to determine the location of each argument on the stack.

## The Stack

**Exercise 9.** *Determine where the kernel initializes its stack, and exactly where in memory its stack is located. How does the kernel reserve space for its stack? And at which "end" of this reserved area is the stack pointer initialized to point to?*

The kernel initializes its stack through the following instructions in `entry.S`

```
movl    $0x0,%ebp          # nuke frame pointer
# Set the stack pointer
movl    $(bootstacktop),%esp
```

`bootstacktop` is defined in `bootasm.S` and is the top of the stack, defined as the end of the bootstack section (defined by `.space` directive, which reserves `KSTKSIZE` bytes of memory for the kernel stack). From `kernel.sym` we can see that the stack starts from the address `0xf0107000`.

`KSTKSIZE` defines the size of the stack. The stack pointer is initialized to point to the end of the bootstack section (which is `bootstacktop`, located at `0xf010f000` in memory). Kernel starts from here in memory and grows downwards (hence, we observe lower memory addresses) in memory. Also, the difference between `0xf010f000` and `0xf0107000` tells the size of the stack which is 8kb.

The kernel reserves the space for its stack using the `.space` directive as follows:

```
.data
#####
# boot stack
#####
    .p2align    PGSHIFT    # force page alignment
    .globl      bootstack
bootstack:
    .space      KSTKSIZE
    .globl      bootstacktop
bootstacktop:
```

**Exercise 10.** How many 32-bit words does each recursive nesting level of `test_backtrace` push on the stack, and what are those words?

The address of the `test_backtrace` function is `0xf0100040`.

```
Breakpoint 1, test_backtrace (x=3) at kern/init.c:13
13      {
(gdb) c
Continuing.
=> 0xf0100040 <test_backtrace>: push    %ebp

Breakpoint 1, test_backtrace (x=2) at kern/init.c:13
13      {
(gdb) p $sp
$5 = (void *) 0xf010ef7c
(gdb) c
Continuing.
=> 0xf0100040 <test_backtrace>: push    %ebp

Breakpoint 1, test_backtrace (x=1) at kern/init.c:13
13      {
(gdb) p $sp
$6 = (void *) 0xf010ef5c
(gdb) □
```

We check the address right after each call of `test_backtrace`. The difference between the addresses of two successive calls is 32 bytes. Therefore, each recursive nesting level of `test_backtrace` pushes 8 words (each word is 4 bytes, hence 32 bytes is 8 words) on the stack.

From the `kernel.asm`, and checking the stack values, it contains the following address.

```
Breakpoint 1, test_backtrace (x=5) at kern/init.c:13
13      {
(gdb) x/8x 0xf0100040
0xf0100040 <test_backtrace>: 0x56e58955    0x0172e853    0xc3810000    0x000102be
0xf0100050 <test_backtrace+16>: 0x8308758b    0x8d5608ec    0xff17d883    0xfae850ff
(gdb) x/8x $esp
0xf010efd0: 0xf01000f4    0x00000005    0x00001aac    0x00000660
0xf010efec: 0x00000000    0x00000000    0x00010094    0x00000000
```

This includes the value of `x(5)`, `x-1(4)`, the return address, previous `ebp`, and some arguments.

```

// Test the stack backtrace function (lab 1 only)
test_backtrace(5);
f01000e8: c7 04 24 05 00 00 00    movl    $0x5, (%esp)
f01000ef: e8 4c ff ff ff          call    f0100040 <test_backtrace>
f01000f4: 83 c4 10                add     $0x10, %esp

// Drop into the kernel monitor.
while (1)
    monitor(NULL);
f01000f7: 83 ec 0c                sub     $0xc, %esp
f01000fa: 6a 00                   push    $0x0
f01000fc: e8 a0 07 00 00          call    f01008a1 <monitor>
f0100101: 83 c4 10                add     $0x10, %esp
f0100104: eb f1                   jmp     f01000f7 <i386_init+0x51>

```

```

// Test the stack backtrace function (lab 1 only)
void
test_backtrace(int x)
{
f0100040: 55                      push    %ebp
f0100041: 89 e5                   mov     %esp, %ebp
f0100043: 56                      push    %esi
f0100044: 53                      push    %ebx
f0100045: e8 72 01 00 00          call    f01001bc <__x86.get_pc_thunk.bx>
f010004a: 81 c3 be 02 01 00        add     $0x102be, %ebx
f0100050: 8b 75 08                mov     0x8(%ebp), %esi
    cprintf("entering test_backtrace %d\n", x);
f0100053: 83 ec 08                sub     $0x8, %esp
f0100056: 56                      push    %esi
f0100057: 8d 83 d8 17 ff ff        lea     -0xe828(%ebx), %eax
f010005d: 50                      push    %eax
f010005e: e8 fa 09 00 00          call    f0100a5d <cprintf>
    if (x > 0)
f0100063: 83 c4 10                add     $0x10, %esp
f0100066: 85 f6                   test    %esi, %esi
f0100068: 7e 29                   jle     f0100093 <test_backtrace+0x53>
        test_backtrace(x-1);
f010006a: 83 ec 0c                sub     $0xc, %esp
f010006d: 8d 46 ff                lea     -0x1(%esi), %eax
f0100070: 50                      push    %eax
f0100071: e8 ca ff ff ff          call    f0100040 <test_backtrace>
f0100076: 83 c4 10                add     $0x10, %esp
}

```

The backtrace code can't detect the number of arguments because the function is not keeping track of the number of arguments passed to it. To fix this, we would need to include debugging symbol table in the kernel, which will provide information about the number and type of arguments for each function. Using this information, the code can then keep track of the number of arguments passed to it.

**Exercise 11.** Implement the backtrace function as specified above.

```
int
mon_backtrace(int argc, char **argv, struct Trapframe *tf)
{
    // Your code here.
    uint32_t *ebp; // base pointer
    asm volatile("movl %%ebp, %0" : "=r" (ebp)); // get base pointer
    // asm volatile is used to prevent compiler from optimizing the code
    uint32_t eip; // instruction pointer
    uint32_t args[5]; // arguments
    int i;

    cprintf("Stack backtrace:\n");
    while(ebp != 0) { // base pointer is 0 at the end of the stack
        eip = *(ebp + 1); // get the instruction pointer
        cprintf("ebp %08x eip %08x args", ebp, eip); // print the
        // base pointer and instruction pointer - %08x is used to print the
        // hexadecimal value of the pointer
        for(i = 0; i < 5; i++) {
            args[i] = *(ebp + 2 + i); // get the arguments
            cprintf(" %08x", args[i]); // print the arguments
        }
        cprintf("\n");
        ebp = (uint32_t *)*ebp; // get the base pointer of the previous
        // stack frame
    }
    return 0;
}
```

Output:

```
Type 'help' for a list of commands.
K> backtrace
Stack backtrace:
bp f010ff58 eip f0100a71 args 00000001 f010ff80 00000000 f0100ad5 f0100a84
    kern/monitor.c:136: monitor+333
bp f010ffd8 eip f0100101 args 00000000 00001aac 00000660 00000000 00000000
    kern/init.c:43: i386_init+91
bp f010fff8 eip f010003e args 00000023 00001003 00002003 00003003 00004003
    kern/entry.S:83: <unknown>+0
K> □
```

(This also includes the output from ex 12 as the screenshot was taken later.)

This can be also implemented using the `read_ebp()` function, as mentioned, as follows:

```
int
mon_backtrace(int argc, char **argv, struct Trapframe *tf)
{
    // Your code here.
    uint32_t *ebp = (uint32_t *)read_ebp(); //base pointer
    uint32_t eip = ebp[1]; //eip is the second element of the stack frame.
    //The first element is the return address of the caller
    cprintf("Stack backtrace:\n");
    while(ebp != 0){
        cprintf("ebp %08x eip %08x args %08x %08x %08x %08x %08x\r\n",
ebp, eip, ebp[2], ebp[3], ebp[4], ebp[5], ebp[6]); //print the stack
frame - ebp, eip, and the first 5 arguments
        ebp = (uint32_t *)*ebp; //move to the next stack frame
        eip = ebp[1]; //update eip - without this, the eip will be the
same as the previous stack frame (we dn't want that)
    }
    return 0;
}
```

**Exercise 12.** Modify your stack backtrace function to display, for each `eip`, the function name, source file name, and line number corresponding to that `eip`.

To print the line number, we add the following code to `kdebug.c`

```
stab_binsearch(stabs, &lline, &rline, N_SLINE, addr); //
N_SLINE is the type for line numbers in stabs - as mentioned in the header
file
    if (lline <= rline) { // if the
line number is found
        info->eip_line = stabs[lline].n_desc; //
n_desc is the line number in the stab and is stored in the stabs structure
- in this line we are storing the line number in the eip_line variable
    } else {
        return -1; // if
the line number is not found, return -1
    }
```



Using the bt function in gdb to verify the implementation of backtrace line number, we can observe that there is some difference in the line numbers

```
Breakpoint 1, test_backtrace (x=5) at kern/init.c:13
13      {
(gdb) bt
#0  test_backtrace (x=5) at kern/init.c:13
#1  0xf01000f4 in i386_init () at kern/init.c:39
#2  0xf010003e in relocated () at kern/entry.S:80
(gdb) █
```

Example – kern/init.c is at 43 in our implementation whereas it is at 30 in gdb. This is because the gdb implementation gives the line number where the function is called.

In `debuginfo_eip`, where do `__STAB_*` come from?

Observe the following from the kern/kernel.ld

```
/* Include debugging information in kernel memory */
.stab : {
    PROVIDE(__STAB_BEGIN__ = .);
    *(.stab);
    PROVIDE(__STAB_END__ = .);
    BYTE(0)      /* Force the linker to allocate space
                  for this section */
}

```

`__STAB_BEGIN__` and `__STAB_END__` are symbols defined in the kernel.ld. These are the beginning and end addresses in the .stab section.

From `objdump -h obj/kern/kernel`, observe the following -

```
2 .stab          00003889 f01022ec 001022ec 000032ec 2**2
                  CONTENTS, ALLOC, LOAD, READONLY, DATA
3 .stabstr       00001629 f0105b75 00105b75 00006b75 2**0
                  CONTENTS, ALLOC, LOAD, READONLY, DATA
```

We can also observe these in kernel.asm

```
// Search the entire set of stabs for the source file (type N_S0).
lfile = 0;
f0100c50: c7 45 e4 00 00 00 00    movl    $0x0, -0x1c(%ebp)
rfile = (stab_end - stabs) - 1;
f0100c57: c7 c0 ec 22 10 f0      mov     $0xf01022ec,%eax
f0100c5d: c7 c2 74 5b 10 f0      mov     $0xf0105b74,%edx
f0100c63: 29 c2                  sub     %eax,%edx
f0100c65: c1 fa 02              sar     $0x2,%edx
f0100c68: 69 d2 ab aa aa aa      imul    $0xaaaaaaaa,%edx,%edx
f0100c6e: 83 ea 01              sub     $0x1,%edx
f0100c71: 89 55 e0              mov     %edx, -0x20(%ebp)
```